

ILCMS Technical Specification v2.0

AI/LLM-Assisted Legal Intelligence, RAG & AI Governance

Icelandic Legal Case Management System (ILCMS)

Developer & Architecture Specification

2 September 2026

Status: Version 2.0 — Engineering Baseline

Predecessor: ILCMS Platform Version 1.0 Developer Handoff

Primary principle: AI assists authorised legal professionals; it does not become the authoritative legal decision-maker.

Document Control and Design Position

- Version 2.0 extends the Version 1 platform with a governed AI subsystem.
- AI is optional and replaceable. Core case management, records, authorization and audit do not depend on an external AI provider.
- RAG is the preferred factual-grounding mechanism for ILCMS content.
- Human approval is mandatory for configured high-impact outputs.
- All AI processing is subject to existing ILCMS authorization and data-classification controls.
- Regulatory statements are engineering guidance and must be validated against current Icelandic/EEA law before a compliance claim is made.

1. Purpose and Scope

- Version 2.0 adds a governed AI/LLM capability layer to the Version 1 ILCMS platform. It covers AI-assisted legal intelligence, retrieval-augmented generation (RAG), document intelligence, drafting assistance, extraction, classification, deadline suggestions, summarisation, semantic search, and AI governance.
- AI is an assistive subsystem. It must not silently become an authoritative source of legal facts, deadlines, case metadata, records, or legal conclusions. Human users remain responsible for legal decisions and approval of high-impact outputs.
- The design supports three deployment modes: fully local/air-gapped, private infrastructure with locally hosted models, and controlled external model providers. Core case management remains operational when all external AI services are disabled.

2. Version 2.0 Objectives

- Provide secure, citation-grounded AI assistance over authorised ILCMS content.
- Prevent cross-case, cross-organisation and privilege-boundary leakage.
- Make every material AI interaction reproducible and auditable.
- Provide model, prompt, retrieval and policy versioning.
- Support human review and explicit approval for high-impact outputs.
- Allow replacement of models and AI infrastructure without redesigning the case-management domain.
- Provide measurable evaluation, quality, security and cost controls.

3. Non-Goals and Safety Boundaries

- ILCMS will not autonomously provide binding legal advice, decide conflicts, determine legal rights, file court documents without explicit human approval, or alter authoritative records based solely on an AI result.
- AI-generated text is draft content until explicitly accepted by an authorised human.
- AI-generated metadata is a proposal until accepted or automatically applied only where the configured risk policy explicitly permits it.
- AI must not be used as the authorization mechanism. Authorization is enforced by the normal ILCMS security model before retrieval and before presentation of results.

4. AI Architecture Principles

- Security boundary first: retrieve only data the requesting principal is already authorised to access.
- Grounding first: prefer authoritative ILCMS documents and structured case data over model memory.
- Traceability first: retain model, prompt/template, retrieval set, source references, policy version and outcome metadata.
- Human-in-the-loop: high-impact operations require explicit approval.
- Provider neutrality: model providers are adapters behind a stable AI service interface.
- Fail closed: if authorization, provenance, model policy, or required citations cannot be established, the operation fails rather than guessing.
- Reproducibility: model and prompt versions are immutable once published; historical AI runs retain their versions.

5. AI Capability Classification

- Class A — Low impact assistance: summarising a selected document, rewriting text, extracting non-authoritative metadata, generating search suggestions.
- Class B — Decision support: case summaries, issue extraction, deadline suggestions, conflict indicators, legal-research assistance, document classification.
- Class C — High impact assistance: proposed legal arguments, proposed court submissions, proposed client communications, recommendations affecting legal strategy, automated external actions.
- Class D — Prohibited autonomous action: autonomous legal decision-making, autonomous filing/submission, autonomous client advice, autonomous alteration/deletion of authoritative records, or any action explicitly prohibited by policy.

6. AI Risk and Governance Framework

- Each AI feature must have a registered AI use case with owner, purpose, data classes, risk class, permitted models, human-review requirement, retention period, evaluation criteria and incident procedure.
- Use-case registration should include whether personal data, privileged material, confidential case material, special-category data, court material or third-party data can enter the workflow.
- Policy decisions are configuration data and are versioned/audited. A deployed AI use case must reference an immutable policy version.
- ILCMS should maintain an AI inventory containing deployed models, providers, versions, capabilities, licences, deployment location, approved use cases and security classification.

7. Regulatory Alignment

- The AI governance design should be capable of supporting applicable Icelandic law, EEA data-protection requirements and, where applicable, EU AI Act controls. This is an engineering baseline, not a legal opinion.
- The EU AI Act is now in its general application phase, with certain provisions having earlier or later application dates. The European Commission states that transparency rules under Article 50 apply from August 2026 and that AI-literacy obligations have applied since February 2025. ILCMS should therefore implement transparency, AI-literacy records, logging, human oversight and risk classification as first-class controls.
- For Iceland, the legal applicability and incorporation status of specific EU/EEA AI instruments must be validated before claiming compliance. The system should therefore label regulatory controls as configurable compliance mappings rather than hard-coded legal assertions.

8. AI Service Logical Architecture

- Frontend → API Gateway → AI Orchestrator → Policy/Authorization Check → Retrieval Service → Prompt Builder → Model Adapter → Output Validator → Human Review/Approval → Case/Document Service.
- Supporting services: Model Registry, Prompt Registry, RAG Index, Embedding Service, AI Audit Store, Evaluation Service, Safety/PII Filter, Job Queue, Usage/Cost Metering and AI Administration UI.
- The AI Orchestrator owns workflow coordination; it does not own case data. Authoritative data remains in PostgreSQL/object storage and normal ILCMS domain services.

9. AI Deployment Topologies

- Air-gapped: Kubernetes-hosted model server, local embedding model, local OCR/extraction, PostgreSQL/pgvector or local search engine, object storage, RabbitMQ and local observability. No internet dependency.
- Private connected: same architecture with controlled access to approved external model endpoints through an AI egress gateway.
- Hybrid: sensitive case content processed locally while non-sensitive tasks may use approved external models. Policy determines routing per use case and data classification.
- External-only AI is not an acceptable mandatory dependency for core ILCMS operation.

10. Model Serving and Model Registry

- Provide a provider-neutral adapter supporting chat/completions, embeddings and structured output.

- Recommended local-serving options may include vLLM or another production-grade inference server; lighter local development may use Ollama or equivalent. The specific runtime is an implementation choice.
- Model registry fields: model_id, name, provider, version, checksum/digest, capability set, context limit, quantisation, licence, deployment location, security classification, approved use cases, evaluation status, active status.
- Model artefacts must be integrity-verified and tracked by immutable digest. Model upgrades require evaluation before production activation.

11. Prompt and Policy Registry

- Prompts are versioned artefacts, not inline strings scattered through application code.
- Prompt record: prompt_id, version, purpose, language, system instructions, variables schema, safety policy, output schema, owner, approval state, created_at and published_at.
- Production prompts are immutable. Changes create a new version.
- Prompt templates must never contain secrets, credentials or hard-coded privileged case data.
- Policy and prompt versions used in an AI run are stored with the run metadata.

12. Retrieval-Augmented Generation (RAG)

- RAG is the default architecture for factual answers over ILCMS content.
- Pipeline: query normalisation → authorization filtering → lexical search → semantic/vector search → metadata filtering → ranking/reranking → source selection → context assembly → generation → citation validation.
- Hybrid retrieval should combine PostgreSQL full-text/pg_trgm and vector search for the initial implementation. The design must allow migration to OpenSearch or another search backend later.
- Every retrieved chunk must retain source document ID, version ID, page/section where available, case ID, classification, access policy and content hash.

13. RAG Authorization Model

- Authorization filtering occurs before model context construction. A document that the user cannot access must never enter the prompt.
- Case-level, document-level, organisation-level, team-level, private-note and legal-hold restrictions must be respected.
- Search indexes and vector stores are derived data. They must contain enough security metadata to support filtering, but the authoritative authorization decision remains the ILCMS authorization service.
- Index rebuilds must reapply current authorization metadata and must not broaden access.

14. Chunking and Document Ingestion

- Chunking must preserve document identity, version, page, heading, paragraph/section boundaries and offsets where feasible.
- Do not blindly split legal documents at arbitrary character counts. Prefer structural boundaries such as headings, paragraphs, numbered clauses and page boundaries.
- Store chunk metadata: chunk_id, document_version_id, case_id, page_start/page_end, section_path, ordinal, content_hash, embedding_model_version, created_at.
- Re-indexing must be idempotent and support model-version migration without destroying the previous index until the replacement is validated.

15. Embeddings and Vector Storage

- Embedding models are registered and versioned exactly like generative models.
- Initial implementation may use PostgreSQL with pgvector to reduce operational complexity. A dedicated vector/search engine can be introduced when scale or retrieval requirements justify it.
- Vectors are derived data and may be deleted/rebuilt. Source documents and authoritative metadata are never dependent on vectors.
- Embedding model changes require full or incremental re-embedding and retrieval-quality evaluation.

16. Citation and Provenance Requirements

- AI answers that make factual claims about case material must provide citations to source documents.
- Citations should identify document name/ID, version, page or section when available, and a stable internal reference.
- The UI should allow the user to open the cited source subject to normal authorization.
- An answer that cannot satisfy required citation policy must be labelled as insufficiently grounded or withheld, depending on use-case risk.
- AI-generated legal propositions should distinguish source-grounded statements from model-generated analysis.

17. Hallucination and Groundedness Controls

- Use retrieval confidence, citation coverage, source relevance, contradiction checks and structured output validation where appropriate.
- Prompts should explicitly instruct the model to state when evidence is unavailable rather than inventing facts.
- For high-risk use cases, require at least one authoritative source citation per material factual proposition.
- Do not treat model confidence scores as factual certainty. They are advisory signals only.
- Evaluation datasets must include adversarial questions, missing-information questions, conflicting documents and deliberately misleading documents.

18. AI Legal Research Assistant

- Provide a case-aware research interface that can search authorised internal documents and, where explicitly configured, approved external legal sources.
- Internal and external sources must be clearly separated.
- External research must identify source URL/provider and retrieval timestamp when allowed.
- AI must not present a generated legal proposition as an authoritative statement of current Icelandic law without a cited source.
- External legal research adapters must be optional so the system remains functional in an air-gapped environment.

19. Case Summarisation

- Support document summary, case summary, chronology, issue summary, party summary, proceeding summary and recent-activity summary.
- Summaries should identify their source scope and generation timestamp.
- Case summaries are derived views, not authoritative case records, unless a user explicitly saves an approved version.
- Approved summaries should retain the generating model/prompt metadata and source snapshot information.

20. Document Intelligence

- AI may propose document type, parties, dates, subjects, classification codes, confidentiality labels and related case links.
- OCR and deterministic metadata extraction should be preferred where sufficient; LLM extraction is a fallback or enrichment layer.
- Extracted fields must carry provenance and confidence/quality metadata.
- Human approval is required for high-impact classification or metadata changes.

21. Deadline and Obligation Intelligence

- AI may identify candidate deadlines, obligations, dates and procedural events from documents.
- AI-derived deadlines enter a PROPOSED state and must not become authoritative deadlines without configured approval.
- Store source document/version, page, extracted text span, model, prompt, confidence and reviewer decision.
- Deterministic deadline rules remain authoritative when a rule can be applied. AI is a discovery/assistance mechanism, not a replacement for configured legal-calendar rules.

22. Conflict-Checking Assistance

- AI may identify aliases, related entities, names and potential relationships that could improve conflict searches.
- AI cannot make the final conflict determination.
- Potential matches must link to underlying authoritative records and require human review under configured policy.
- Conflict-assistance operations must receive elevated audit treatment because false negatives can have serious consequences.

23. Drafting Assistant

- Provide controlled drafting for internal notes, correspondence, summaries, chronologies, document outlines and proposed legal text.
- Drafts must be visibly labelled as AI-assisted until accepted.
- External/client/court-facing communications require explicit human approval.
- Approved drafts become normal ILCMS documents/records through the standard document lifecycle; the AI draft itself is retained only according to configured policy.

24. AI Chat / Copilot UX

- AI chat should be case-scoped by default. The user must explicitly choose broader scopes.
- Show active context sources, data scope, model, and applicable AI-use policy in the UI.
- Allow users to inspect citations, remove selected documents from context, regenerate, provide feedback and report an unsafe/incorrect result.
- Never expose hidden system prompts or confidential policy internals to end users, but provide meaningful transparency about how the answer was produced.

25. AI API Design

- Use `/api/v1/ai/...` for AI capabilities.
- Suggested endpoints: POST `/ai/runs`, GET `/ai/runs/{id}`, POST `/ai/chat`, POST `/ai/summarize`, POST `/ai/extract`, POST `/ai/research`, POST `/ai/deadline-suggestions`, POST `/ai/draft`, GET `/ai/sources/{id}`.
- Long-running operations return 202 with job/run identifiers.
- Every AI mutation supports Idempotency-Key where retries can occur.
- AI responses should include `run_id`, `status`, `model_version`, `prompt_version`, `policy_version`, `citations` and `generated content` where applicable.

26. AI Run Data Model

- `ai_runs`: `id`, `organization_id`, `case_id`, `user_id`, `use_case`, `risk_class`, `status`, `model_id`, `model_version`, `prompt_id`, `prompt_version`, `policy_version`, `input_scope`, `started_at`, `completed_at`, `token/input-output metrics`, `error_code`, `approved_at`, `approved_by`.
- `ai_run_sources`: `run_id`, `document_version_id`, `chunk_id`, `retrieval_rank`, `retrieval_score`, `citation_key`.
- `ai_run_outputs`: `run_id`, `output_type`, `content`, `structured_payload`, `validation_status`, `created_at`.
- `ai_feedback`: `run_id`, `user_id`, `rating`, `category`, `comment`, `created_at`.
- `ai_incidents`: `id`, `run_id`, `category`, `severity`, `description`, `status`, `owner`, `resolution`, `timestamps`.

27. AI Audit and Provenance

- Audit every production AI run at a level appropriate to the use case.
- Minimum audit fields: `who`, `when`, `case scope`, `use case`, `model/version`, `prompt/version`, `policy/version`, `retrieval source IDs`, `output status`, `approval/rejection`, and `relevant system request ID`.
- Sensitive prompt content should not automatically be duplicated into ordinary logs. Store controlled AI run records with explicit retention.
- Audit records must be tamper-evident and subject to the same operational controls as Version 1 audit events.

28. Data Protection and Confidentiality

- AI processing must respect the same data-classification and access controls as the rest of ILCMS.
- Do not send case data to an external provider unless the organisation has explicitly enabled that provider for the relevant data class and use case.
- Provider contracts, processing locations, retention behaviour and training/data-use policies must be represented in provider metadata and governance controls.
- Support data minimisation, purpose limitation, retention limits, deletion workflows and data-subject rights processes as applicable.
- The EDPB's Opinion 28/2024 highlights data-protection questions around AI-model development and deployment, including legal basis and anonymisation; these concerns should inform ILCMS governance and DPIA processes.

29. AI Egress Gateway

- All external AI traffic should pass through a controlled egress service.
- Enforce provider allowlists, destination restrictions, TLS, authentication, rate limits, payload policy, data-classification policy and audit logging.
- Block external transmission when policy cannot establish that the data class is permitted.
- Air-gapped installations omit the external egress path entirely.

30. AI Security Controls

- Threat model prompt injection, indirect prompt injection, malicious documents, data exfiltration, model supply-chain compromise, insecure tools, excessive agency, retrieval poisoning and denial-of-service.
- Treat document text as untrusted input. A document must never be able to override system policy merely because it contains imperative instructions.
- Separate system instructions, retrieved content and user instructions.
- Tool use must be allowlisted and scoped. AI must never receive unrestricted shell, database, network or administrative capabilities.
- Use sandboxing for risky parsers and document-processing workloads.

31. AI Tool/Agent Model

- Version 2 should initially use a non-agentic orchestration model: bounded tools, explicit workflows and deterministic permissions.
- If agents are introduced later, every tool must declare input schema, authorization requirement, side effects, risk class, confirmation requirement and audit event.
- High-impact tools must require explicit human confirmation immediately before execution.
- AI cannot grant itself permissions or broaden retrieval scope.

32. Evaluation Framework

- Maintain versioned evaluation datasets for summarisation, retrieval, citation accuracy, extraction, deadline detection, classification and drafting.
- Measure retrieval recall/precision, citation correctness, groundedness, factual accuracy, extraction F1 where applicable, latency, cost and refusal behaviour.
- Evaluate in Icelandic and English, including Icelandic legal terminology, inflection, names, dates and document conventions.
- Every model/prompt/index change must pass regression evaluation before production release.
- Production feedback should feed a controlled evaluation pipeline rather than silently changing prompts.

33. AI Red-Team and Adversarial Testing

- Test cross-case leakage, prompt injection, malicious PDFs, hidden text, poisoned metadata, indirect instructions, jailbreaks, citation spoofing, conflicting evidence and oversized context.
- Test authorization at retrieval, prompt assembly, response display and export boundaries.
- Test that disabling AI providers does not break core case management.

- Maintain security regression tests for every discovered AI vulnerability.

34. AI Observability

- Metrics: AI runs, success/failure, latency, queue depth, tokens, model utilisation, GPU utilisation where applicable, retrieval latency, citation coverage, safety blocks, external-provider calls and estimated cost.
- Logs must be structured and privacy-aware. Do not log full prompts or confidential document content by default.
- Distributed tracing should correlate API request → AI run → retrieval → model call → output validation → approval.
- Alert on abnormal external data transmission, sudden retrieval failures, model error rates, latency spikes and unusual usage patterns.

35. AI Resource and Cost Governance

- Track usage by organisation, case, user, model and use case.
- Support quotas and rate limits.
- Separate compute costs for local inference from provider/API costs.
- Use smaller models for low-risk tasks and larger models only where justified by evaluation results.
- Cost controls must never weaken security or authorization.

36. Kubernetes / Infrastructure Design

- Deploy AI components as independent workloads: ai-orchestrator, model-server, embedding-server, retrieval service, evaluation worker and optional GPU node pool.
- Use node labels/taints for GPU scheduling, resource requests/limits, PodDisruptionBudgets, NetworkPolicies and secrets from Vault.
- Model artefacts should be pulled from an approved internal registry/object store in restricted environments.
- GPU workloads require dedicated observability for memory, utilisation, queueing and thermal/health signals where available.

37. AI Data Lifecycle

- Source document → extraction/OCR → chunking → embedding → index → retrieval → AI run → output → human approval → optional authoritative record.
- Each stage has separate provenance and retention rules.
- Derived AI artefacts must be rebuildable.
- Deleting or placing a document under legal hold must trigger the appropriate index and retrieval lifecycle operations.

38. AI Governance Administration

- Admin screens: AI providers, models, model versions, prompts, policies, use cases, risk classifications, quotas, approved data classes, external providers, evaluation suites, incidents and audit reports.
- Production configuration requires appropriate administrative privilege and should support approval for high-impact changes.
- Model and prompt publication should follow Draft → Evaluation → Approved → Published → Retired.
- Retired versions remain available for historical audit but cannot be selected for new production runs.

39. Human Oversight and Approval

- Approval UI must show the AI output, sources, relevant warnings, model/prompt/policy version and proposed side effects.
- Approval actions: ACCEPT, ACCEPT_WITH_EDITS, REJECT, ESCALATE.
- Approval records identify the human approver and timestamp.
- Approval must not retroactively imply that the AI output was correct; it records a human decision to use the result.

40. AI Incident Management

- Incident categories: privacy breach, authorization leakage, hallucination, unsafe recommendation, prompt injection, provider failure, model defect, data poisoning, incorrect deadline, incorrect citation and policy violation.
- Each incident captures severity, affected runs/cases, containment, investigation, remediation and closure.
- Critical incidents can disable a model, provider, prompt or use case through an emergency kill switch.
- Incident actions are audited.

41. Testing Strategy

- Unit tests for policy evaluation, authorization filtering, chunking, citation mapping and structured output validation.
- Integration tests for model adapters, vector retrieval, PostgreSQL/pgvector, object storage and job queues.
- Security tests for tenant isolation and prompt-injection scenarios.
- End-to-end tests for representative legal workflows.
- Offline tests must be possible without external model providers.

42. Implementation Phases

- P0 foundation: AI domain model, provider abstraction, model registry, policy registry, audit model, local model adapter, embedding adapter and feature flags.
- P1 RAG MVP: document chunking, pgvector, hybrid retrieval, authorization filtering, citations and case-scoped chat.
- P2 document intelligence: summarisation, extraction, classification proposals and approved metadata workflows.
- P3 legal intelligence: deadline suggestions, chronology, issue extraction, internal legal research and drafting assistance.
- P4 governance maturity: evaluation platform, red-team suite, incident management, quotas, model lifecycle automation and advanced observability.
- P5 optional advanced agents: only after bounded tool governance and human-approval controls are mature.

43. Initial P0/P1 Developer Tickets

- AI-001 Create AI domain models and migrations.
- AI-002 Implement Model Registry and provider adapter interface.
- AI-003 Implement Prompt Registry and versioning.
- AI-004 Implement AI Policy Registry and risk classes.
- AI-005 Implement AI Run/Audit persistence.
- AI-006 Implement local embedding adapter.
- AI-007 Add pgvector and document chunk tables.
- AI-008 Implement authorization-aware hybrid retrieval.
- AI-009 Implement citation model and source resolver.
- AI-010 Implement case-scoped RAG chat endpoint.
- AI-011 Implement local model adapter.
- AI-012 Implement AI feature flags and kill switch.
- AI-013 Add AI observability metrics/tracing.
- AI-014 Build adversarial retrieval/prompt-injection test suite.
- AI-015 Build evaluation dataset and regression harness.

44. Definition of Done for AI Features

- Threat model completed.
- Data classification and permitted scope documented.
- Authorization tests pass.

- Prompt/model/policy versions are registered.
- AI run is auditable.
- Sources/citations are available where required.
- Human approval is implemented for the configured risk class.
- Offline/local execution has been tested where the feature is required to support air-gapped operation.
- Evaluation results meet the feature's acceptance thresholds.
- Operational metrics, alerts and rollback/disable procedures exist.
- User-facing AI transparency is present.

45. Version 2.0 Acceptance Criteria

- Core ILCMS continues to operate with AI disabled.
- No AI request can retrieve content outside the caller's existing authorization scope.
- Every production AI run is attributable to a user/service identity and versioned model/prompt/policy.
- Case-grounded answers provide verifiable citations according to policy.
- High-impact AI outputs cannot become authoritative or externally transmitted without configured human approval.
- Models and prompts can be replaced without changing the case-management domain.
- RAG indexes can be rebuilt from authoritative data.
- AI provider egress can be disabled centrally.
- Air-gapped deployment can execute the supported local AI feature set.
- AI evaluation and security regression tests are part of CI/CD before production release.

46. Recommended Version 2.0 Technology Baseline

- Application: existing FastAPI/Python 3.12+ backend and worker architecture.
- Database: PostgreSQL + pgvector initially.
- Search: PostgreSQL FTS/pg_trgm + vector retrieval; abstraction retained for OpenSearch.
- Object storage: existing S3-compatible storage.
- Messaging: existing RabbitMQ.
- Cache: existing Redis.
- Local inference: vLLM or equivalent production model server.
- Embedding/reranking: locally hosted models selected through the model registry.
- Secrets: existing Vault/External Secrets pattern.
- Observability: existing OpenTelemetry/Prometheus/Grafana/logging stack.
- Kubernetes: dedicated optional GPU node pool with NetworkPolicies and resource controls.

47. Final Design Principle

- ILCMS Version 2.0 should be an AI-assisted legal information system, not an autonomous legal actor. The system's competitive advantage should come from combining authoritative case records, secure retrieval, Icelandic legal-document intelligence, excellent provenance, and disciplined human oversight.
- The architecture deliberately makes AI replaceable. Models will change faster than the legal domain and the case-management core. Therefore PostgreSQL/object storage, authorization, records management, audit, workflow and legal-process data remain authoritative; AI remains a governed interpretation and assistance layer over that foundation.

Appendix A — Core AI Data Model Summary

Entity	Purpose	Key controls
ai_models	Registered model/provider versions	digest, licence, approval, deployment location
ai_prompts	Versioned prompt templates	immutable published versions
ai_policies	Use-case/risk/data rules	versioned and audited
ai_runs	Execution provenance	identity, model, prompt, policy, sources
ai_run_sources	RAG provenance	document/version/chunk/rank
ai_outputs	Generated results	validation, approval state
ai_feedback	User evaluation	rating/category/audit
ai_incidents	AI incidents	severity, containment, remediation

Appendix B — Key API Resource Inventory

- /api/v1/ai/runs
- /api/v1/ai/chat
- /api/v1/ai/summarize
- /api/v1/ai/extract
- /api/v1/ai/research
- /api/v1/ai/deadline-suggestions
- /api/v1/ai/draft
- /api/v1/ai/sources/{id}
- /api/v1/admin/ai/providers
- /api/v1/admin/ai/models
- /api/v1/admin/ai/prompts
- /api/v1/admin/ai/policies
- /api/v1/admin/ai/evaluations
- /api/v1/admin/ai/incidents

Appendix C — Non-Negotiable AI Engineering Rules

1. Never use AI as an authorization mechanism.
2. Never place unauthorized content into an AI prompt.
3. Never treat a generated answer as an authoritative legal record without explicit human action.
4. Never permit an AI tool to grant itself permissions.
5. Never make external AI connectivity mandatory for core ILCMS operation.
6. Never silently change model, prompt or policy versions for a reproducible production workflow.
7. Never log confidential prompts/document content by default.
8. Never allow AI-generated deadlines to bypass the authoritative deadline workflow.
9. Never treat vector indexes as authoritative data.
10. Never deploy a new production model without regression evaluation.
11. Always provide provenance appropriate to the risk class.
12. Always retain an emergency kill switch for models/providers/use cases.

Appendix D — Regulatory/Reference Notes

The design intentionally aligns with widely used AI governance principles and should be mapped to the legal obligations actually applicable to the deployment. The European Commission states that the EU AI Act's transparency rules apply from August 2026 and describes risk-based obligations, while the EDPB Opinion 28/2024 addresses data-protection issues arising in AI-model development and deployment. Iceland-specific applicability should be confirmed separately before making a compliance assertion.

Reference sources: European Commission AI Act materials; European Commission AI literacy and transparency guidance; European Data Protection Board Opinion 28/2024.